

Caffe

Tian Huang

2016011379

April 11, 2017

t-huang16@mails.tsinghua.edu.cn

Installation

- ◆ Mac OS 10.12 Sierra
- ◆ Homebrew (package manager)
- ◆ CUDA 7 (useless without GPU)
- ◆ Anaconda Python
- ◆ Modify environment variables

Installation

- ◆ install OpenCV & hdf5

```
brew install -vd snappy leveldb gflags glog szip lmdb  
# need the homebrew science source for OpenCV and hdf5  
brew tap homebrew/science  
brew install hdf5 opencv
```

- ◆ if using Anaconda python, hdf5 can be skipped.
- ◆ brew edit opencv

```
args << "-DPYTHON_LIBRARY=#{py_prefix}/lib/libpython2.7.dylib"  
args << "-DPYTHON_INCLUDE_DIR=#{py_prefix}/include/python2.7"
```


Installation

- ◆ with Python
 - ◆ `brew install --build-from-source --with-python -vd protobuf`
 - ◆ `brew install --build-from-source -vd boost boost-python`
- ◆ without Python
 - ◆ `brew install protobuf boost`

Installation

- ◆ BLAS (or MKL)
- ◆ Python (without Anaconda)
- ◆ for req in \$(cat requirements.txt); do pip install \$req; done

Compilation

```
cp Makefile.config.example Makefile.config
# Adjust Makefile.config (for example, if using Anaconda Python, or if cuDNN is
desired)
make all
make test
make runtest
```

- ◆ Notice:
- ◆ Without GPU:
- ◆ uncomment CPU_ONLY := 1 in Makefile.config

```
ANACONDA_HOME := $(HOME)/anaconda
PYTHON_INCLUDE := $(ANACONDA_HOME)/include \
                 $(ANACONDA_HOME)/include/python2.7 \
                 $(ANACONDA_HOME)/lib/python2.7/site-packages/numpy/core/includ
e \
```


Compilation

- ◆ if you forget it

```
[tinahtdeMacBook-Pro:caffe-master Tinaht1$ make all
NVCC src/caffe/layers/absval_layer.cu
nvcc fatal   : Unsupported gpu architecture 'compute_60'
make: *** [.build_release/cuda/src/caffe/layers/absval_layer.o] Error 1
```

- ◆ when you try to modify

```
caffe/caffe/device/device.cu in common.o
ld: symbol(s) not found for architecture x86_64
clang: error: linker command failed with exit code 1 (use -v to see invocation)
make: *** [.build_release/lib/libcaffe.so.1.0.0-rc5] Error 1
[tinahtdeMacBook-Pro:caffe-master Tinaht1$
```

- ◆ maybe no way out

Compilation

- ◆ to compile python or matlab
- ◆ make pycaffe
- ◆ make matcaffe
- ◆ set paths in Makefile.config before make

Compilation

- ◆ python 2.7 is ok, python 3.6 isn't?

```
[>>> import caffe
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  File "//Users/Tinaht1/Desktop/大一春/oop/caffe-master/python/caffe/__init__.py", line 1, in <module>
    from .pycaffe import Net, SGDSolver, NesterovSolver, AdaGradSolver, RMSPropSolver, AdaDeltaSolver, AdamSolver, NCCL, Timer
  File "//Users/Tinaht1/Desktop/大一春/oop/caffe-master/python/caffe/pycaffe.py", line 13, in <module>
    from ._caffe import Net, SGDSolver, NesterovSolver, AdaGradSolver, \
ImportError: dynamic module does not define module export function (PyInit__caffe)
[>>> quit()]
```


Compilation

```
[-----] 4 tests from SoftmaxWithLossLayerTest/0, where TypeParam = N5caffe9CPUDeviceIfEE
[ RUN      ] SoftmaxWithLossLayerTest/0.TestForwardIgnoreLabel
[          OK ] SoftmaxWithLossLayerTest/0.TestForwardIgnoreLabel (1 ms)
[ RUN      ] SoftmaxWithLossLayerTest/0.TestGradient
[          OK ] SoftmaxWithLossLayerTest/0.TestGradient (5 ms)
[ RUN      ] SoftmaxWithLossLayerTest/0.TestGradientUnnormalized
[          OK ] SoftmaxWithLossLayerTest/0.TestGradientUnnormalized (5 ms)
[ RUN      ] SoftmaxWithLossLayerTest/0.TestGradientIgnoreLabel
[          OK ] SoftmaxWithLossLayerTest/0.TestGradientIgnoreLabel (6 ms)
[-----] 4 tests from SoftmaxWithLossLayerTest/0 (17 ms total)

[-----] 1 test from MultinomialLogisticLossLayerTest/1, where TypeParam = d
[ RUN      ] MultinomialLogisticLossLayerTest/1.TestGradientCPU
[          OK ] MultinomialLogisticLossLayerTest/1.TestGradientCPU (0 ms)
[-----] 1 test from MultinomialLogisticLossLayerTest/1 (0 ms total)

[-----] 2 tests from HDF5DataLayerTest/0, where TypeParam = N5caffe9CPUDeviceIfEE
[ RUN      ] HDF5DataLayerTest/0.TestRead
[          OK ] HDF5DataLayerTest/0.TestRead (4 ms)
[ RUN      ] HDF5DataLayerTest/0.TestSkip
[          OK ] HDF5DataLayerTest/0.TestSkip (20 ms)
[-----] 2 tests from HDF5DataLayerTest/0 (24 ms total)


[-----] Global test environment tear-down
[=====] 1104 tests from 150 test cases ran. (48058 ms total)
[ PASSED  ] 1104 tests.
```


cmake

- ◆ if you fail to make it, use cmake
- ◆ if it doesn't work either,
- ◆ keep calm, uninstall them and go back to the first page.


Deep Learning

- ◆ NN training each layer's parameter can be replaced by its last layer's " $\sigma(w*x+b)$ "



Deep Learning

the first layer's w^*x+b



the last layer's $h(x)$

$$h(x) = \sigma(\sum w_i x_i + b) = \sigma(w \cdot x)$$

Training LeNet on MNIST



Training LeNet on MNIST

Prepare Datasets

- ◆ under "caffe_root"
- ◆ `./data/mnist/get_mnist.sh`
- ◆ `./example/mnist/create_mnist.sh`
- ◆ maybe wget isn't installed: `brew install wget`

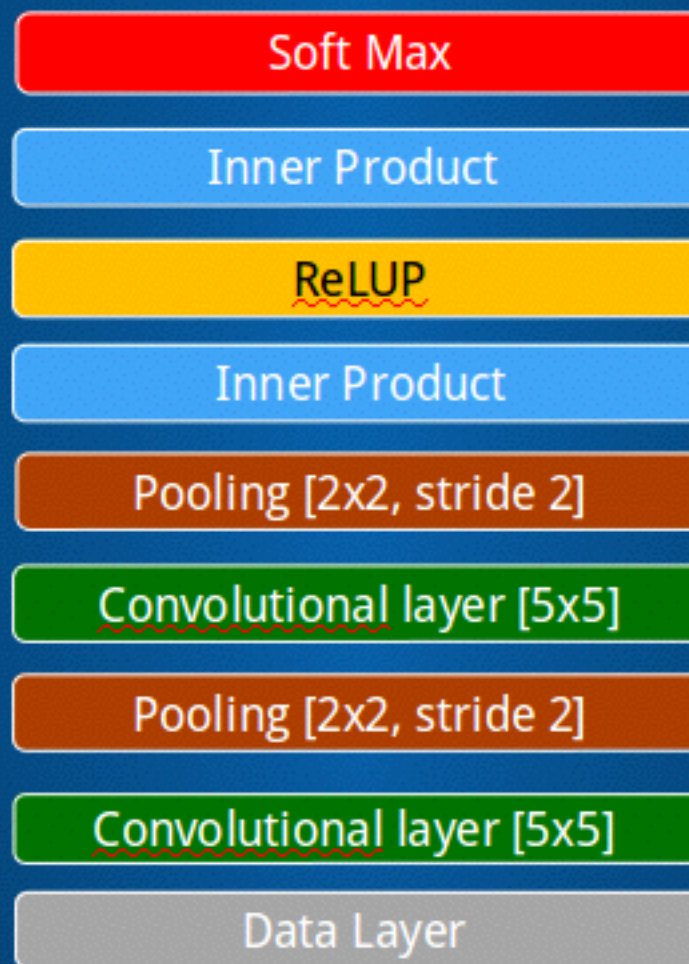
Define the MNIST Network

- ◆ define by protobuf :
- ◆ `src/caffe/proto/caffe.proto`
- ◆ Seven types of Layers (defined by prototxt)
- ◆ `/examples/mnist/lenet_train_test.prototxt`
- ◆ Each layer consists of two parts:
- ◆ type & parameter

Layers

LeNet topology

FORWARD



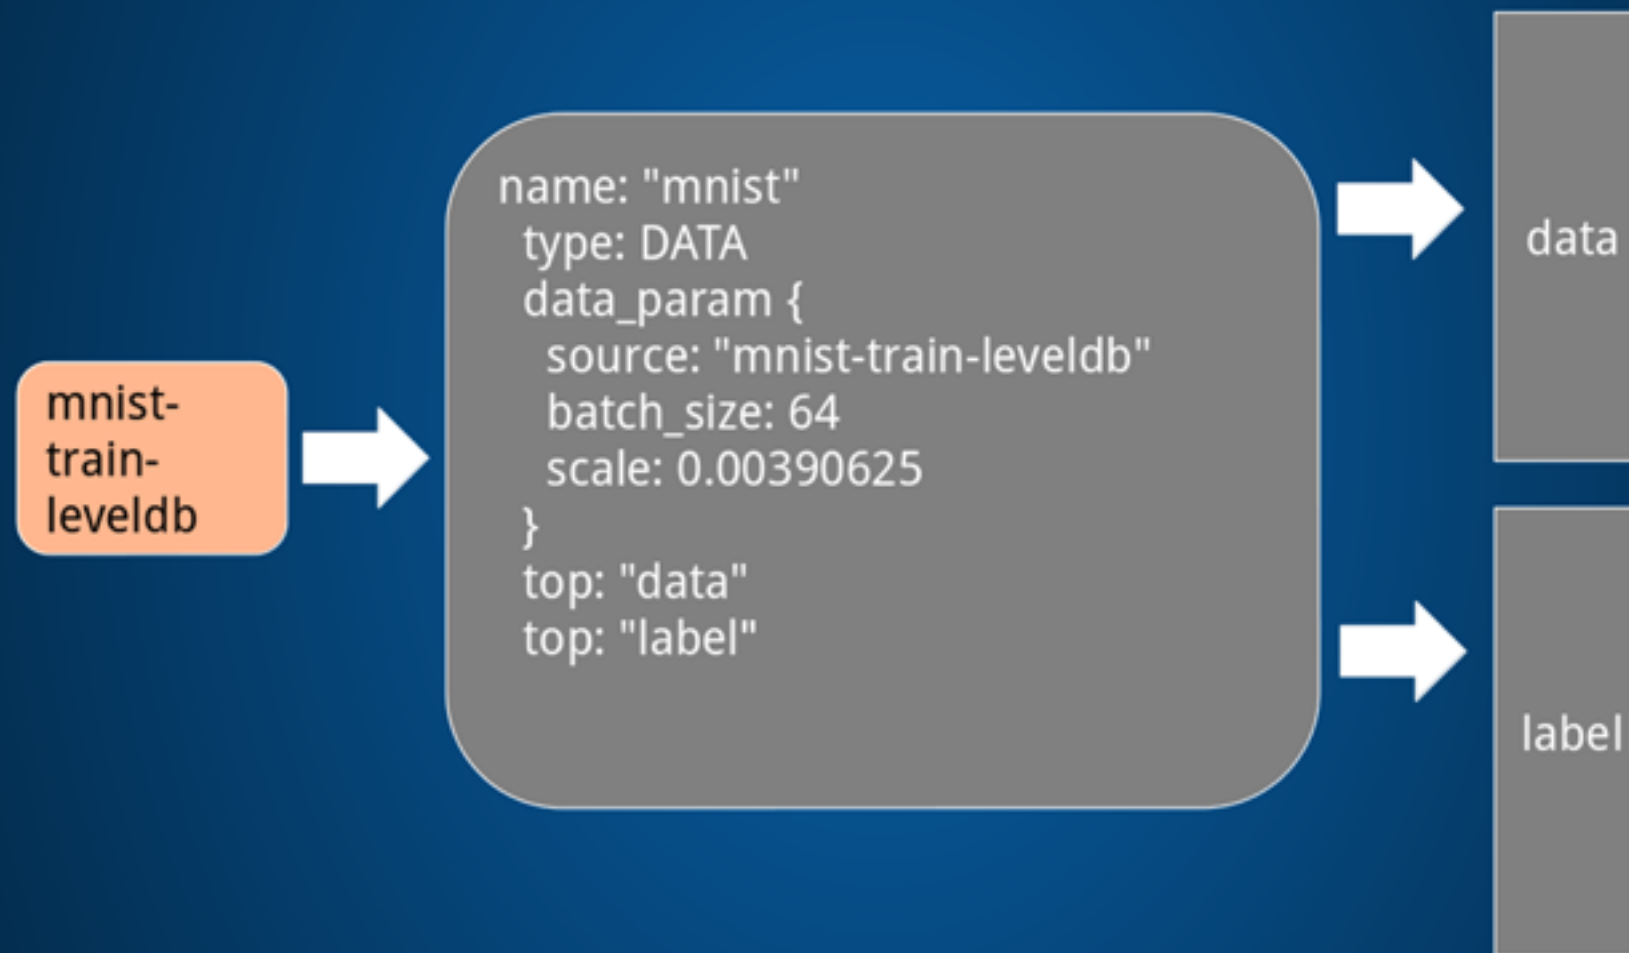
BACKWARD

The Data Layer

```
2 layer {
3   name: "mnist"
4   type: "Data"
5   top: "data"
6   top: "label"
7   include {
8     phase: TRAIN ->usable only in "train" phase
9   }
10  transform_param {
11    scale: 0.00390625 ->Image pixel value, *1/256
12  }
13  data_param {
14    source: "examples/mnist/mnist_train_lmdb"
15    batch_size: 64
16    backend: LMDB
17  }
18 }
```


The Data Layer

Data Layer

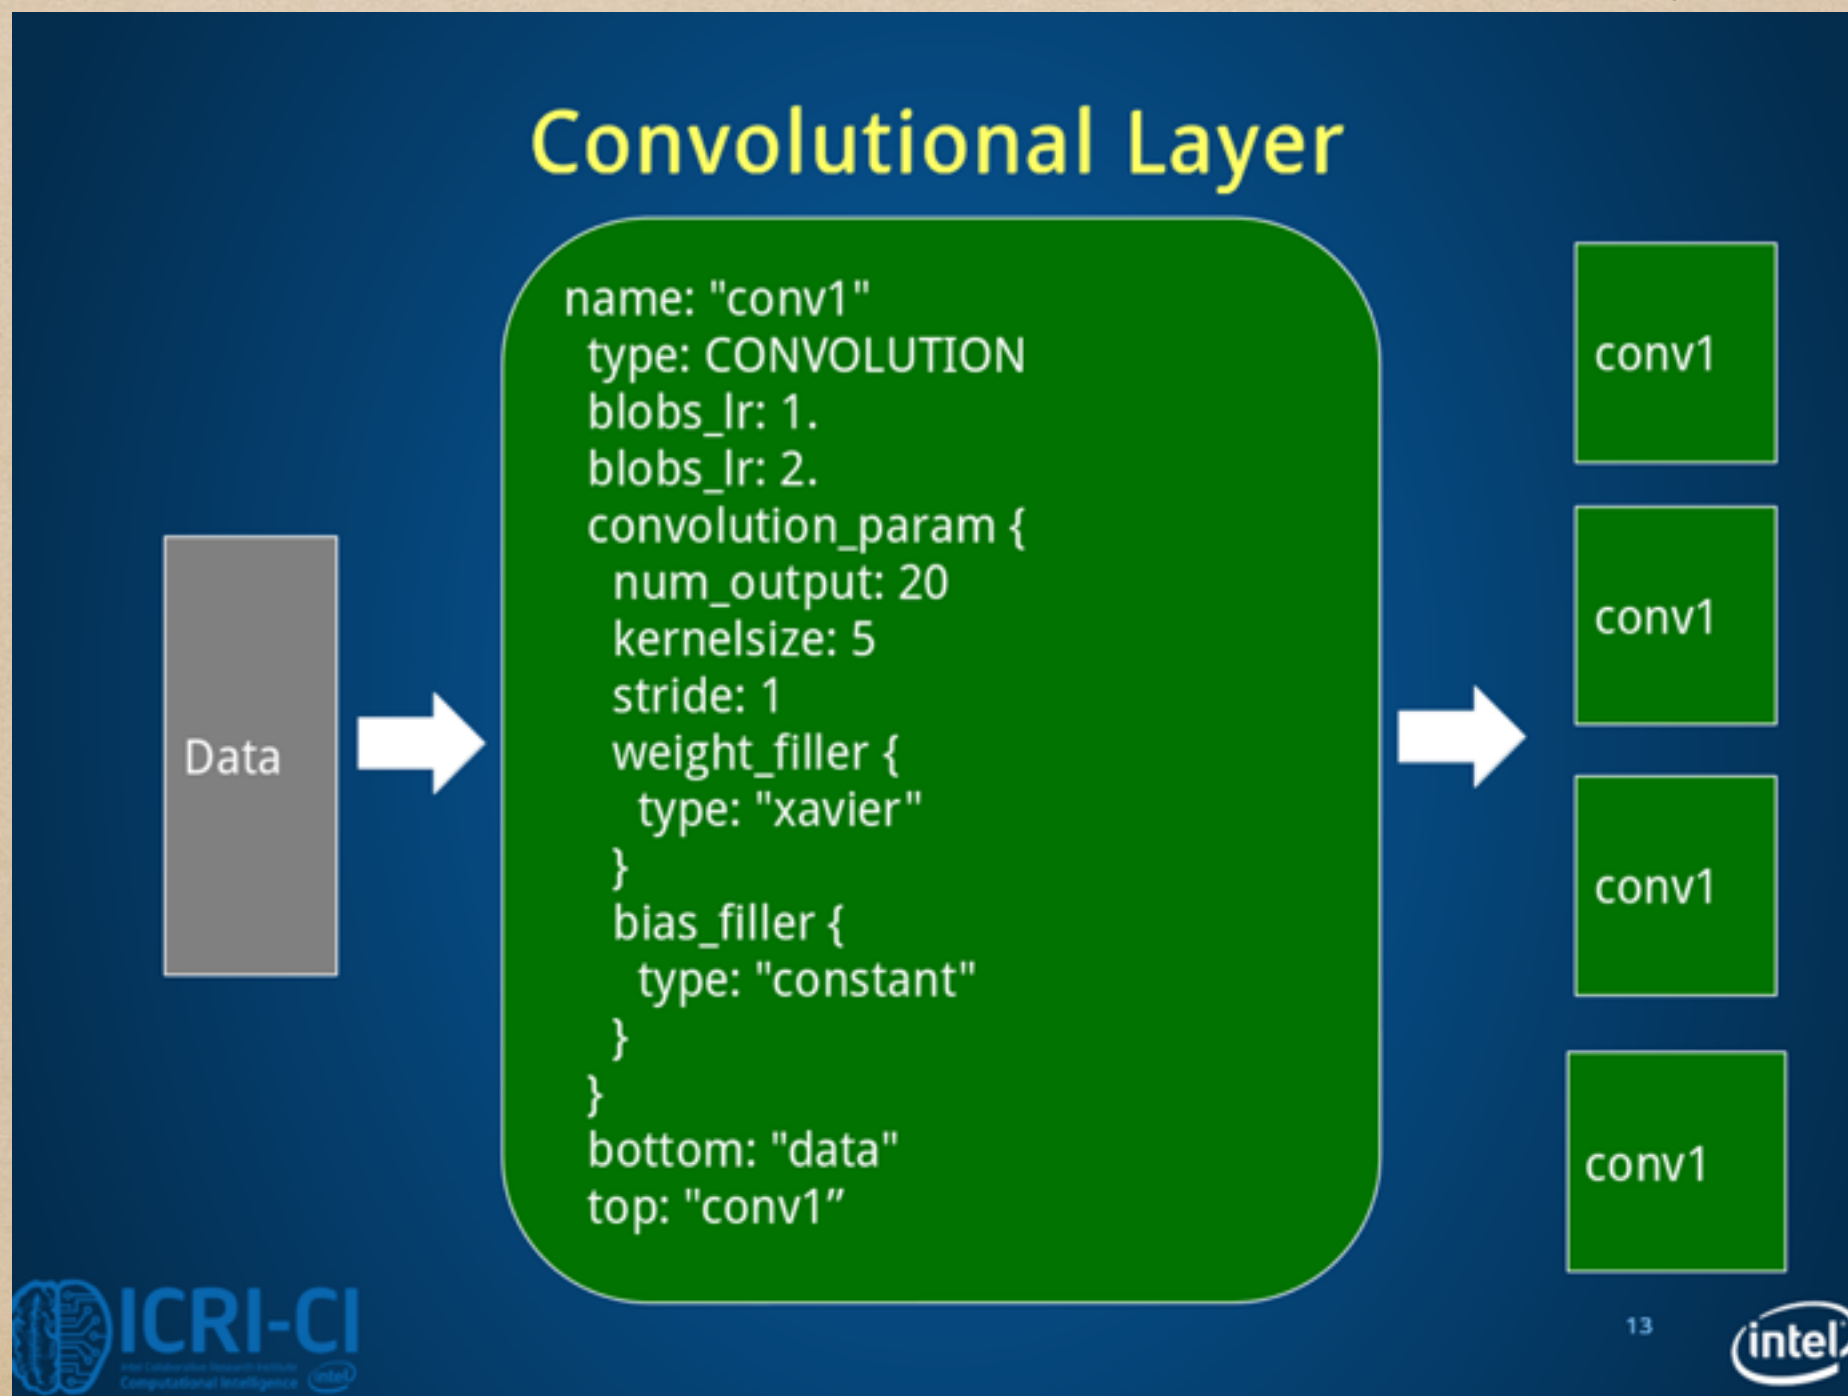


The Convolution Layer

```
36 layer {
37     name: "conv1"
38     type: "Convolution"
39     bottom: "data" ->input
40     top: "conv1" ->output
41     param {
42         lr_mult: 1 ->times of w learning rate
43     }
44     param {
45         lr_mult: 2 ->times of b learning rate
46     }
47     convolution_param {
48         num_output: 20
49         kernel_size: 5
50         stride: 1
51         weight_filler {
52             type: "xavier" ->initialize w
53         }
54         bias_filler {
55             type: "constant" ->initialize b, const
56         }
57     }
58 }
```


$$a_1 = \sigma(b + w * a_0)$$

The Convolution Layer




The Convolution Layer


input neurons




hidden neuron



The Convolution Layer



The Convolution Layer



the Pooling Layer

```
59 layer {  
60     name: "pool1"  
61     type: "Pooling"  
62     bottom: "conv1"  
63     top: "pool1"  
64     pooling_param {  
65         pool: MAX  
66         kernel_size: 2  
67         stride: 2  
68     }  
69 }
```

L2 pooling
max pooling

->no overlapping

the Pooling Layer

Pooling Layer


```
name: "pool1"  
type: POOLING  
pooling_param {  
  kernel_size: 2  
  stride: 2  
  pool: MAX  
}  
bottom: "conv1"  
top: "pool1"
```



```
for (p = 0; p < k; p++)  
  for (q = 0; q < k; q++)  
     $y_L(x, y) = \max(y_L(x, y), y_{L-1}(x*s + p, y*s + q));$ 
```

Poolinh helps to extract features that are increasingly invariant to local transformations of the input image.

the Pooling Layer



the Fully Connected Layer

```
104 layer {
105     name: "ip1"
106     type: "InnerProduct"
107     bottom: "pool2"
108     top: "ip1"
109     param {
110         lr_mult: 1
111     }
112     param {
113         lr_mult: 2
114     }
115     inner_product_param {
116         num_output: 500
117         weight_filler {
118             type: "xavier"
119         }
120         bias_filler {
121             type: "constant"
122         }
123     }
124 }
```

change multiple
dimensions into
 $N*1*1$

the Fully Connected Layer

Inner product (Fully Connected) Layer

```
name: "ip1"  
type: INNER_PRODUCT  
blobs_lr: 1.  
blobs_lr: 2.  
inner_product_param {  
  num_output: 500  
  weight_filler {  
    type: "xavier"  
  }  
  bias_filler {  
    type: "constant"  
  }  
}  
bottom: "pool2"  
top: "ip1"
```

$$Y_L(n) = \sum W_L(n, m) * Y_{L-1}(m)$$

The ReLU Layer

```
125 layer {  
126     name: "relu1"  
127     type: "ReLU"  
128     bottom: "ip1"  
129     top: "ip1"  
130 }
```

the same name to the bottom and top blobs
in-place operations to save some memory

After the ReLU layer, another innerproduct layer

Add:

negative_slope

The ReLU Layer

ReLU Layer

```
layers {  
  name: "relu1"  
  type: RELU  
  bottom: "ip1"  
  top: "ip1"  
}
```

$$Y_L(n; x, y) = \max(Y_{L-1}(n; x, y), 0);$$

$$h(x) = \sigma(\sum w_i x_i + b) = \sigma(w x)$$

PReLU?

```
type: "PReLU"
```

```
...
```

```
prelu_param{
```

```
  filler: {
```

```
    value: 0.25 (default)
```

```
  }
```

```
  channel_shared: false
```

```
}
```


the Accuracy Layer

```
152 layer {  
153     name: "accuracy"  
154     type: "Accuracy"  
155     bottom: "ip2"  
156     bottom: "label"  
157     top: "accuracy"  
158     include {  
159         phase: TEST ->usable only in "test" phase  
160     }  
161 }
```

report the model accuracy

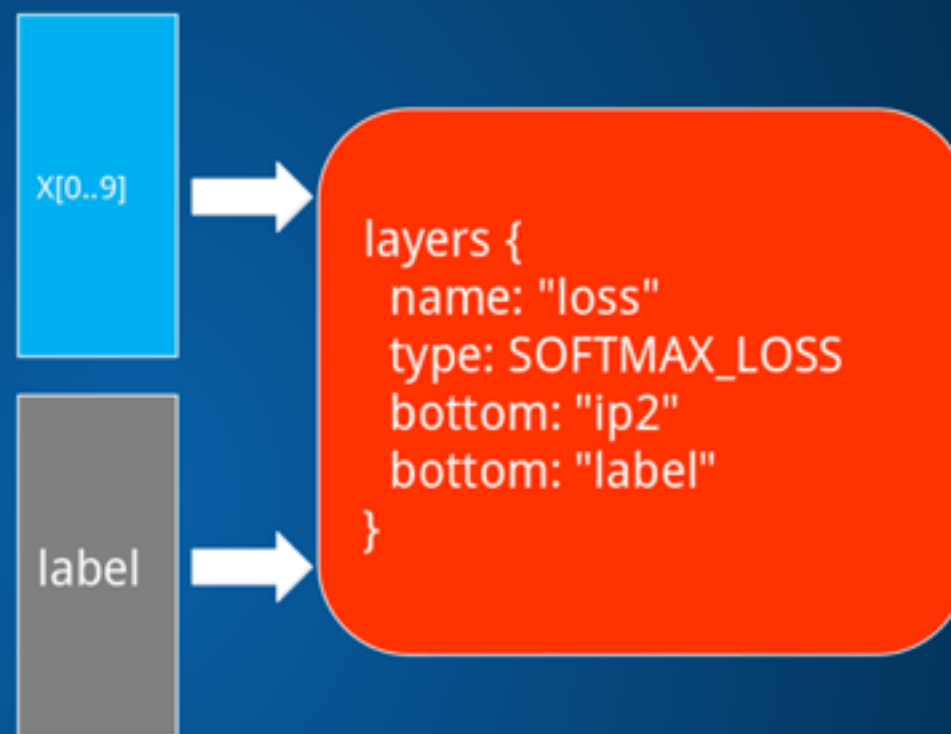
the Softmax Layer

```
162 layer {  
163     name: "loss"  
164     type: "SoftmaxWithLoss"  
165     bottom: "ip2"  
166     bottom: "label"  
167     top: "loss"  
168 }
```

to compute the loss function value
two blobs: ip2 (prediction)
 & label (provided by data)

the Loss Layer

SoftMax + Loss Layer



Combines softmax:

$$Y_L[i] = \exp(Y_{L-1}[i]) / (\sum(Y_L[i]));$$

with log-loss :

$$E = -\log(Y_L(\text{label}(n)))$$

Define the MNIST Solver

/examples/mnist/lenet_solver.prototxt

```
1 # The train/test net protocol buffer definition
2 net: "examples/mnist/lenet_train_test.prototxt"
3 # test_iter specifies how many forward passes the test should carry
  out.
4 # In the case of MNIST, we have test batch size 100 and 100 test
  iterations,
5 # covering the full 10,000 testing images.
6 test_iter: 100
7 # Carry out testing every 500 training iterations.
8 test_interval: 500
9 # The base learning rate, momentum and the weight decay of the
  network.
10 base_lr: 0.01
11 momentum: 0.9
12 weight_decay: 0.0005
13 # The learning rate policy
14 lr_policy: "inv"
15 gamma: 0.0001
16 power: 0.75
17 # Display every 100 iterations
18 display: 100
19 # The maximum number of iterations
20 max_iter: 10000
21 # snapshot intermediate results
22 snapshot: 5000
23 snapshot_prefix: "examples/mnist/lenet"
24 # solver mode: CPU or GPU
25 solver_mode: CPU
26
```


Training and Testing the Model

`./examples/mnist/train_lenet.sh`

```
I0411 15:12:40.709614 249827328 data_layer.cpp:73] Restarting data prefetching from start.  
I0411 15:12:40.824187 3886519232 solver.cpp:398] Test net output #0: accuracy = 0.9901  
I0411 15:12:40.824270 3886519232 solver.cpp:398] Test net output #1: loss = 0.0278956 (* 1 = 0.0278956 loss)  
I0411 15:12:40.875422 3886519232 solver.cpp:219] Iteration 9000 (14.0311 iter/s, 7.127s/100 iters), loss = 0.0170253  
I0411 15:12:40.875460 3886519232 solver.cpp:238] Train net output #0: loss = 0.0170251 (* 1 = 0.0170251 loss)  
I0411 15:12:40.875473 3886519232 sgd_solver.cpp:105] Iteration 9000, lr = 0.00617924  
I0411 15:12:44.790310 3886519232 solver.cpp:219] Iteration 9100 (25.5493 iter/s, 3.914s/100 iters), loss = 0.00823545  
I0411 15:12:44.790349 3886519232 solver.cpp:238] Train net output #0: loss = 0.00823528 (* 1 = 0.00823528 loss)
```


```
I0411 15:13:34.475559 249827328 data_layer.cpp:73] Restarting data prefetching from start.  
I0411 15:13:34.603086 3886519232 solver.cpp:398] Test net output #0: accuracy = 0.9908  
I0411 15:13:34.603121 3886519232 solver.cpp:398] Test net output #1: loss = 0.0266228 (* 1 = 0.0266228 loss)
```

CPU:

Time: about 10 min (10000 iterations)

Accuracy: 99%

Cifar 10



Picture Recognition

train: 50000

test: 10000

32*32

10 types

Cifar 10

- ◆ database:
- ◆ `./data/cifar10/get_cifar10.sh`
- ◆ `./examples/cifar10/create_cifar10.sh`
- ◆ a little problem (`create_cifar10.sh` should be run just under “caffe root”)

Training

- ◆ `./examples/cifar10/train_quick.sh`
- ◆ Get:
- ◆ `cifar10_quick_iter_4000.caffemodel.h5` (model)
- ◆ `cifar10_quick_iter_4000.solverstate.h5` (state)

Result

```
I0411 16:51:20.566542 103510016 data_layer.cpp:73] Restarting data prefetching from start.  
I0411 16:51:21.023499 3886519232 solver.cpp:398] Test net output #0: accuracy = 0.715  
I0411 16:51:21.023531 3886519232 solver.cpp:398] Test net output #1: loss = 0.842517 (* 1 = 0.842517 loss)  
I0411 16:51:21.023537 3886519232 solver.cpp:316] Optimization Done.  
I0411 16:51:21.023541 3886519232 caffe.cpp:259] Optimization Done.
```

- ◆ really slowly....
- ◆ CPU only:
- ◆ Time: about 20min (4000 iterations)
- ◆ Accuracy: 70%
- ◆ I Need a GPU....

using caffe training model

- ◆ 1、prepare datasets
- ◆ 2、write *_test.prototxt for modeling
- ◆ 3、write *_solver.prototxt for optimization
- ◆ 4、run it in the command line
- ◆ Then you'll get a training model.

Thank you~